

Name :- Sandesh A. Mohite

Roll NO :- 23

Div:- A

Batch:- TA1

Example: A Simple Student Dataset

[]

```
import pandas as pd

# A tiny hand-made dataset to illustrate the concepts above
student_data = {
    "RollNo": [1, 2, 3, 4, 5],
    "Name": ["Aarav", "Diya", "Kabir", "Meera", "Ishan"],
    "Age": [19, 20, 19, 21, 20],
    "Gender": ["M", "F", "M", "F", "M"],
    "CGPA": [8.4, 9.1, 7.8, 8.9, 7.2],
    "Placed": ["Yes", "Yes", "No", "Yes", "No"],
}

student_df = pd.DataFrame(student_data)
student_df # displaying the DataFrame
```

	RollNo	Name	Age	Gender	CGPA	Placed
0	1	Aarav	19	M	8.4	Yes
1	2	Diya	20	F	9.1	Yes
2	3	Kabir	19	M	7.8	No
3	4	Meera	21	F	8.9	Yes
4	5	Ishan	20	M	7.2	No

2. Import Libraries

2. Import Libraries

Library	Why we need it
NumPy	Fast numerical computing; the foundation every other library is built on
Pandas	Loading, cleaning, and manipulating tabular data (DataFrames)
Matplotlib	The base plotting library — full control over every visual element
Seaborn	Beautiful statistical plots built on top of Matplotlib, works great with Pandas
Scikit-learn	Preprocessing tools (scaling, encoding) and ML algorithms

```
# Core libraries
import numpy as np
import pandas as pd

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Preprocessing (used later for scaling & encoding)
from sklearn.preprocessing import (
    StandardScaler, MinMaxScaler, RobustScaler, Normalizer,
    LabelEncoder, OneHotEncoder, OrdinalEncoder
)

# Display / plotting settings
%matplotlib inline
sns.set_theme(style="whitegrid")
plt.rcParams["figure.figsize"] = (7, 4)
pd.set_option("display.max_columns", None)
```

```
print("All libraries imported successfully!")

All libraries imported successfully!
```

3. Loading Datasets

3. Loading Datasets

We'll load **five** classic, beginner-friendly datasets. Each teaches a different EDA skill.

Dataset	Load with	Best for
Titanic	<code>sns.load_dataset('titanic')</code>	Missing values, categorical data, full cleaning
Iris	<code>sns.load_dataset('iris')</code>	Pure numerical analysis, pairplots
Tips	<code>sns.load_dataset('tips')</code>	GroupBy, aggregation, visualization
Penguins	<code>sns.load_dataset('penguins')</code>	Missing values, mixed data types
California Housing	<code>sklearn.datasets.fetch_california_housing</code>	Scaling & correlation on larger numeric data

Titanic will be our main "storyline" dataset — we'll keep coming back to it throughout the notebook. The others appear where they best illustrate a concept.

```
# Load all datasets we will use in this notebook
titanic = sns.load_dataset("titanic")
iris = sns.load_dataset("iris")
tips = sns.load_dataset("tips")
penguins = sns.load_dataset("penguins")

print("Titanic :", titanic.shape)
print("Iris :", iris.shape)
print("Tips :", tips.shape)
print("Penguins:", penguins.shape)
```

```
Titanic : (891, 15)
Iris : (150, 5)
Tips : (244, 7)
Penguins: (344, 7)
```

```
# California Housing – a larger, purely numerical dataset (fetched from scikit-learn)
from sklearn.datasets import fetch_california_housing

try:
    housing_bunch = fetch_california_housing(as_frame=True)
    housing = housing_bunch.frame
except Exception as e:
    # Fallback: if there is no internet access, build a similar synthetic dataset
    print("Could not download California Housing, generating a similar synthetic dataset instead.")
    from sklearn.datasets import make_regression
    X, y = make_regression(n_samples=2000, n_features=8, noise=15, random_state=42)
    cols = ["MedInc", "HouseAge", "AveRooms", "AveBedrms",
            "Population", "AveOccup", "Latitude", "Longitude"]
    housing = pd.DataFrame(X, columns=cols)
    housing["MedHouseVal"] = y

print("California Housing:", housing.shape)
housing.head()
```

California Housing: (20640, 9)

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedHouseVal
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

Check Built-in Dataset from a library

```
import seaborn as sns
sns.get_dataset_names()
```

```
['anagrams',
 'anscombe',
 'attention',
 'brain_networks',
 'car_crashes',
 'diamonds',
 'dots',
 'dowjones',
 'exercise',
 'flights',
 'fmri',
 'geyser',
 'glue',
 'healthexp',
 'iris',
 'mpg',
 'penguins',
 'planets',
 'seaice',
 'taxis',
 'tips',
 'titanic']
```

The `dir()` function lists all attributes, functions, classes, and datasets available in a module.

```
import seaborn as sns
dir(sns)
```

```
['FacetGrid',
 'JointGrid',
 'PairGrid',
 '__builtins__',
 '__cached__',
 '__doc__',
 '__file__',
 '__loader__',
 '__name__',
 '__package__',
 '__path__',
 '__spec__',
 '__version__',
 '_base',
 '_compat',
 '_core',
 '_docstrings',
 '_orig_rc_params',
 '_statistics',
 '_stats',
 '_algorithms',
 '_axes_style',
 '_axisgrid',
 '_barplot',
 '_blend_palette',
 '_boxenplot',
 '_boxplot',
 '_categorical',
 '_catplot',
 '_choose_colorbrewer_palette',
 '_choose_cubehelix_palette',
 '_choose_dark_palette',
 '_choose_diverging_palette',
 '_choose_light_palette',
 '_clustermap',
 '_cm',
 '_color_palette',
 '_colors',
 '_countplot',
 '_crayon_palette',
 '_crayons',
 '_cubehelix_palette',
 '_dark_palette',
 '_desaturate',
 '_despine',
 '_displot',
 '_distplot',
 '_distributions',
 '_diverging_palette',
 '_dogplot',
```

```
'ecdfplot',
'external',
'get_data_home',
'get_dataset_names',
'heatmap',
'histplot',
'hls_palette',
```

4. First Look at a Dataset

Before doing anything else, always run these commands on a new dataset (using Titanic).

```
titanic.head() # first 5 rows – a quick peek at the data
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

```
titanic.tail() # last 5 rows – useful to check if data ends correctly (no junk rows)
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
886	0	2	male	27.0	0	0	13.00	S	Second	man	True	NaN	Southampton	no	True
887	1	1	female	19.0	0	0	30.00	S	First	woman	False	B	Southampton	yes	True
888	0	3	female	NaN	1	2	23.45	S	Third	woman	False	NaN	Southampton	no	False
889	1	1	male	26.0	0	0	30.00	C	First	man	True	C	Cherbourg	yes	True
890	0	3	male	32.0	0	0	7.75	Q	Third	man	True	NaN	Queenstown	no	True

Interpretation: tail() shows the last 5 rows. Checking both ends helps catch issues like footer text or corrupted rows accidentally included when a CSV was exported.

```
titanic.sample(5) # 5 random rows – helps you see rows head()/tail() might miss
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
748	0	1	male	19.0	1	0	53.100	S	First	man	True	D	Southampton	no	False
167	0	3	female	45.0	1	4	27.900	S	Third	woman	False	NaN	Southampton	no	False
567	0	3	female	29.0	0	4	21.075	S	Third	woman	False	NaN	Southampton	no	False
336	0	1	male	29.0	1	0	66.600	S	First	man	True	C	Southampton	no	False
760	0	3	male	NaN	0	0	14.500	S	Third	man	True	NaN	Southampton	no	True

```
print("Shape (rows, columns):", titanic.shape)
print("Size (total values) :", titanic.size)
print("Dimensions (ndim) :", titanic.ndim)
```

```
Shape (rows, columns): (891, 15)
Size (total values) : 13365
Dimensions (ndim) : 2
```


Interpretation: Titanic has **891 rows and 15 columns** → `size` = $891 \times 15 = 13,365$ total data points. `ndim` = 2 because a DataFrame is always 2-dimensional (rows $\times$ columns).

```
print("Columns:\n", list(titanic.columns))
print("\nIndex:", titanic.index)
```

```
Columns:
['survived', 'pclass', 'sex', 'age', 'sibsp', 'parch', 'fare', 'embarked', 'class', 'who', 'adult_male', 'deck', 'embark_town',
Index: RangeIndex(start=0, stop=891, step=1)
```

5. Understanding Variable Types

5. Understanding Variable Types



Type	Definition	Titanic Example
Continuous	Can take any value in a range	<code>fare</code> , <code>age</code>
Discrete	Countable whole numbers	<code>sibsp</code> (siblings/spouses), <code>parch</code>
Nominal	Categories with no order	<code>sex</code> , <code>embarked</code>
Ordinal	Categories with order	<code>pclass</code> (1st > 2nd > 3rd class)
Binary	Exactly two categories	<code>survived</code> (0/1), <code>alone</code> (True/False)

Tip: Correctly identifying variable type decides *which plot* and *which preprocessing method* (scaling vs encoding) you'll use later in this notebook.

```
# Let's classify Titanic's columns ourselves
titanic.dtypes
```

```

0
survived    int64
pclass      int64
sex         object
age         float64
sibsp       int64
parch       int64
fare        float64
embarked    object
class       category
who         object
adult_male  bool
deck        category
embark_town object
alive       object
alone       bool
dtype: object
```

6. Dataset Summary

```
titanic.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   survived    891 non-null    int64
1   pclass      891 non-null    int64
2   sex         891 non-null    object
3   age         714 non-null    float64
4   sibsp       891 non-null    int64
5   parch       891 non-null    int64
6   fare        891 non-null    float64
7   embarked    889 non-null    object
8   class       891 non-null    category
9   who         891 non-null    object
10  adult_male  891 non-null    bool
11  deck        203 non-null    category
12  embark_town 889 non-null    object
13  alive       891 non-null    object
14  alone       891 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
```

```
titanic.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   survived    891 non-null    int64
1   pclass      891 non-null    int64
2   sex         891 non-null    object
3   age         714 non-null    float64
4   sibsp       891 non-null    int64
5   parch       891 non-null    int64
6   fare        891 non-null    float64
7   embarked    889 non-null    object
8   class       891 non-null    category
9   who         891 non-null    object
10  adult_male  891 non-null    bool
11  deck        203 non-null    category
12  embark_town 889 non-null    object
13  alive       891 non-null    object
14  alone       891 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
```


titanic.describe(include="all") # summary for ALL columns (numeric + categorical)

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embar
count	891.000000	891.000000	891	714.000000	891.000000	891.000000	891.000000	889	891	891	891	203	
unique	NaN	NaN	2	NaN	NaN	NaN	NaN	3	3	3	2	7	
top	NaN	NaN	male	NaN	NaN	NaN	NaN	S	Third	man	True	C	South
freq	NaN	NaN	577	NaN	NaN	NaN	NaN	644	491	537	537	59	
mean	0.383838	2.308642	NaN	29.699118	0.523008	0.381594	32.204208	NaN	NaN	NaN	NaN	NaN	
std	0.486592	0.836071	NaN	14.526497	1.102743	0.806057	49.693429	NaN	NaN	NaN	NaN	NaN	
min	0.000000	1.000000	NaN	0.420000	0.000000	0.000000	0.000000	NaN	NaN	NaN	NaN	NaN	
25%	0.000000	2.000000	NaN	20.125000	0.000000	0.000000	7.910400	NaN	NaN	NaN	NaN	NaN	
50%	0.000000	3.000000	NaN	28.000000	0.000000	0.000000	14.454200	NaN	NaN	NaN	NaN	NaN	
75%	1.000000	3.000000	NaN	38.000000	1.000000	0.000000	31.000000	NaN	NaN	NaN	NaN	NaN	
max	1.000000	3.000000	NaN	80.000000	8.000000	6.000000	512.329200	NaN	NaN	NaN	NaN	NaN	

7. Statistical Concepts

7. Statistical Concepts

Let's understand the numbers behind `describe()` using the `age` column.



Concept	Meaning	Formula / Note
Mean	Average value	sum of values ÷ count
Median	Middle value when sorted	robust to outliers
Mode	Most frequently occurring value	works for numeric & categorical
Variance	Average squared distance from the mean	spread of data
Std. Deviation	Square root of variance	spread, in original units
Quartiles	Values that split sorted data into 4 equal parts	Q1, Q2 (median), Q3
Percentile	Value below which a given % of data falls	e.g. 90th percentile
Range	Max - Min	total spread
IQR	Q3 - Q1	spread of the middle 50%, used for outlier detection

```
age = titanic["age"].dropna() # drop missing values just for this calculation
```

8. Missing Values

8. Missing Values

How do missing values happen?

- A survey question left blank (e.g. someone didn't report their `age`).
- Sensor/equipment failure while recording data.
- Data lost during merging of multiple sources.
- Intentionally not applicable (e.g. `deck` unknown for many Titanic passengers).

```
missing_count = titanic.isnull().sum()
missing_percent = (titanic.isnull().sum() / len(titanic) * 100).round(2)

missing_summary = pd.DataFrame({
    "Missing Count": missing_count,
    "Missing %": missing_percent
}).sort_values("Missing %", ascending=False)

missing_summary[missing_summary["Missing Count"] > 0]
```

	Missing Count	Missing %
deck	688	77.22
age	177	19.87
embarked	2	0.22
embark_town	2	0.22

9. Duplicate Values

```
print("Number of duplicate rows in Titanic:", titanic.duplicated().sum())

# Let's create an artificial example to see duplicates in action
demo = titanic.head(5).copy()
demo = pd.concat([demo, demo.iloc[[0, 2]]], ignore_index=True) # manually duplicate 2 rows
print("\nDemo dataset shape (with duplicates):", demo.shape)
print("Duplicate rows found:", demo.duplicated().sum())
```

```
demo[demo.duplicated(keep=False)] # show all copies of duplicated rows
```

Number of duplicate rows in Titanic: 107

Demo dataset shape (with duplicates): (7, 15)

Duplicate rows found: 2

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.250	S	Third	man	True	NaN	Southampton	no	False
2	1	3	female	26.0	0	0	7.925	S	Third	woman	False	NaN	Southampton	yes	True
5	0	3	male	22.0	1	0	7.250	S	Third	man	True	NaN	Southampton	no	False
6	1	3	female	26.0	0	0	7.925	S	Third	woman	False	NaN	Southampton	yes	True

```
demo_cleaned = demo.drop_duplicates()
print("Shape before:", demo.shape, " | Shape after drop_duplicates():", demo_cleaned.shape)
```

Shape before: (7, 15) | Shape after drop_duplicates(): (5, 15)

10. Unique Values

```
print("Unique classes in 'pclass':", titanic["pclass"].unique())
print("Number of unique classes :", titanic["pclass"].nunique())

print("\nValue counts for 'embarked':")
print(titanic["embarked"].value_counts())

print("\nValue counts (as %) for 'sex':")
print(titanic["sex"].value_counts(normalize=True).round(3) * 100)
```

Unique classes in 'pclass': [3 1 2]
Number of unique classes : 3

Value counts for 'embarked':
embarked
S 644
C 168
Q 77
Name: count, dtype: int64

Value counts (as %) for 'sex':
sex
male 64.8
female 35.2
Name: proportion, dtype: float64

Interpretation:

Interpretation:


- `unique()` → lists the *distinct* values in a column.
- `nunique()` → counts how many distinct values exist (great for spotting ID-like columns or columns with only 1 unique value, both of which are useless for modelling).
- `value_counts()` → frequency of each category — instantly reveals **class imbalance** (e.g. far more passengers embarked from 'S' than 'Q').

Tip: If `nunique()` equals the number of rows, that column is likely an identifier (like `PassengerId`) and should usually be dropped before modelling.

11. Correlation

✓ 11. Correlation

Correlation measures how strongly two numerical variables move together, from **-1 to +1**.




- **Positive correlation** → both variables increase together (e.g. `fare` & `pclass` reversed, since 1st class = expensive).
- **Negative correlation** → one increases as the other decreases.
- **Zero / near-zero correlation** → no linear relationship.

Correlation ≠ Causation! Two variables can correlate strongly without one causing the other.

```
numeric_titanic = titanic.select_dtypes(include=np.number)
corr_matrix = numeric_titanic.corr()
corr_matrix
```


	survived	pclass	age	sibsp	parch	fare
survived	1.000000	-0.338481	-0.077221	-0.035322	0.081629	0.257307
pclass	-0.338481	1.000000	-0.369226	0.083081	0.018443	-0.549500
age	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	0.096067
sibsp	-0.035322	0.083081	-0.308247	1.000000	0.414838	0.159651
parch	0.081629	0.018443	-0.189119	0.414838	1.000000	0.216225
fare	0.257307	-0.549500	0.096067	0.159651	0.216225	1.000000

```
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", center=0)
plt.title("Correlation Heatmap — Titanic (numeric columns)")
plt.show()
```



```
plt.figure(figsize=(7, 4))
sns.boxplot(x=titanic["fare"], color="orange")
plt.title("Boxplot of Fare — Outliers appear as individual dots")
plt.show()
```

Boxplot of Fare — Outliers appear as individual dots



Method 2 — IQR Rule (mathematical)

```
Q1 = titanic["fare"].quantile(0.25)
Q3 = titanic["fare"].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

outliers = titanic[(titanic["fare"] < lower_bound) | (titanic["fare"] > upper_bound)]
print(f"Lower bound: {lower_bound:.2f}    Upper bound: {upper_bound:.2f}")
print("Number of outliers in 'fare':", outliers.shape[0])
```

```
Lower bound: -26.72    Upper bound: 65.63
Number of outliers in 'fare': 116
```

Method 3 — Z-score (introduction)

```
from scipy import stats

fare_no_na = titanic["fare"].dropna()
z_scores = np.abs(stats.zscore(fare_no_na))

z_outliers = fare_no_na[z_scores > 3]    # points more than 3 std-devs from the mean
print("Number of outliers using Z-score (|z| > 3):", z_outliers.shape[0])
```

```
Number of outliers using Z-score (|z| > 3): 20
```

Removing Outliers

```
titanic_no_outliers = titanic[(titanic["fare"] >= lower_bound) & (titanic["fare"] <= upper_bound)]
print("Original shape :", titanic.shape)
print("After removing fare outliers:", titanic_no_outliers.shape)
```

```
Original shape : (891, 15)
After removing fare outliers: (775, 15)
```

13. Feature Scaling

13. Feature Scaling

Why scale data?

Many ML algorithms (KNN, SVM, K-Means, Gradient Descent-based models, Neural Networks) are **distance-based** or **gradient-based**. If one feature ranges 0–1 and another 0–100,000, the large-range feature will dominate the distance calculation / gradient — even if it isn't actually more important!

WITHOUT SCALING	WITH SCALING
Income: 0 ————— 200000	Income: -1 — 0 — +1

Age: 0 — 100

(Income completely dominates)

Age: -1 — 0 — +1

(Both features contribute fairly)


```

housing_num = housing.select_dtypes(include=np.number).dropna()
sample_cols = housing_num.columns[:2] # take 2 numeric columns to compare visually

scalers = {
    "Original": housing_num[sample_cols],
    "StandardScaler": pd.DataFrame(StandardScaler().fit_transform(housing_num[sample_cols]), columns=sample_cols),
    "MinMaxScaler": pd.DataFrame(MinMaxScaler().fit_transform(housing_num[sample_cols]), columns=sample_cols),
    "RobustScaler": pd.DataFrame(RobustScaler().fit_transform(housing_num[sample_cols]), columns=sample_cols),
    "Normalizer": pd.DataFrame(Normalizer().fit_transform(housing_num[sample_cols]), columns=sample_cols),
}

fig, axes = plt.subplots(1, 5, figsize=(22, 4))
for ax, (name, data) in zip(axes, scalers.items()):
    sns.histplot(data[sample_cols[0]], ax=ax, kde=True, color="teal")
    ax.set_title(name)
plt.suptitle(f"Effect of Different Scalers on '{sample_cols[0]}'")
plt.tight_layout()
plt.show()

```



Tip: Tree-based models (Decision Tree, Random Forest, XGBoost) do **not** require feature scaling — they split on thresholds, not distances.

Quick Check

- Why would `MinMaxScaler` be risky on a column that has extreme outliers?

14. Encoding Categorical Data

ML algorithms need numbers, not text — encoding converts categories into numeric form.

```

demo_enc = titanic[["sex", "embarked", "pclass"]].dropna().copy()

# 1. Label Encoding — assigns each category an integer (0, 1, 2, ...)
le = LabelEncoder()
demo_enc["sex_label"] = le.fit_transform(demo_enc["sex"])

# 2. One-Hot Encoding — creates a separate 0/1 column for every category
one_hot = pd.get_dummies(demo_enc["embarked"], prefix="embarked")

# 3. Ordinal Encoding — like Label Encoding, but you CONTROL the order
# (pclass 1 = best, 3 = worst, so we keep that natural order)
ordinal = OrdinalEncoder(categories=[[1, 2, 3]])
demo_enc["pclass_ordinal"] = ordinal.fit_transform(demo_enc[["pclass"]])

result = pd.concat([demo_enc[["sex", "sex_label"]], one_hot,
                    demo_enc[["pclass", "pclass_ordinal"]]], axis=1)
result.head()


```

	sex	sex_label	embarked_C	embarked_Q	embarked_S	pclass	pclass_ordinal
0	male	1	False	False	True	3	2.0
1	female	0	True	False	False	1	0.0
2	female	0	False	False	True	3	2.0
3	female	0	False	False	True	1	0.0
4	male	1	False	False	True	3	2.0

15. Data Visualization — The Complete Toolkit This is the biggest and most important section. For each plot type you'll see: Purpose → When to use → Code → Interpretation. We'll mix Titanic, Iris, and Tips datasets to show each plot on the data it suits best.

1. Histogram Purpose: Shows the distribution (shape/spread) of a single numerical variable. When to use: Understanding how one numeric column is spread out — skewed? normal? bimodal?

```
plt.figure(figsize=(7, 4))
sns.histplot(titanic["age"], bins=30, kde=True, color="steelblue")
plt.title("Histogram of Passenger Age (Titanic)")
plt.xlabel("Age")
plt.show()
```




Interpretation: Most passengers are between 20–40 years old; there's a small secondary bump near very young children — the distribution is slightly right-skewed (a long tail toward older ages)

2. Countplot

Purpose: Shows the *frequency* of each category in a categorical column. **When to use:** Comparing counts across categories (class imbalance checks)

```
plt.figure(figsize=(6, 4))
sns.countplot(x="pclass", data=titanic, hue="pclass", palette="Set2", legend=False)
plt.title("Passenger Count by Class")
plt.show()
```




Interpretation: 3rd class had by far the most passengers, followed by 1st, then 2nd — the dataset is imbalanced across classes.

3. Barplot

Purpose: Shows the *average* (or another statistic) of a numeric variable across categories. **When to use:** Comparing a summary statistic (like mean survival rate) between groups.

```
plt.figure(figsize=(6, 4))
sns.barplot(x="pclass", y="survived", data=titanic, hue="pclass", palette="Set1", legend=False)
plt.title("Average Survival Rate by Class")
plt.ylabel("Survival Rate")
plt.show()
```



4. Boxplot

Purpose: Shows median, quartiles, and outliers of a numeric variable — optionally split by category. **When to use:** Comparing distributions & spotting outliers across groups


```
plt.figure(figsize=(7, 4))
sns.boxplot(x="pclass", y="age", data=titanic, hue="pclass", palette="pastel", legend=False)
plt.title("Age Distribution by Passenger Class")
plt.show()
```



5. Violin Plot

Purpose: Combines a boxplot with a KDE (density) curve, showing the full *shape* of the distribution. **When to use:** When you want more distributional detail than a boxplot alone gives.

```
plt.figure(figsize=(7, 4))
sns.violinplot(x="sex", y="age", data=titanic, hue="sex", palette="muted", legend=False)
plt.title("Age Distribution by Gender (Violin Plot)")
plt.show()
```




```
plt.figure(figsize=(7, 4)) sns.violinplot(x="sex", y="age", data=titanic, hue="sex", palette="muted", legend=False) plt.title("Age Distribution by Gender (Violin Plot)") plt.show()
```

6. Scatter Plot

Purpose: Shows the relationship between **two** numerical variables. **When to use:** Spotting correlation, trends, or clusters between two continuous variables.


```
plt.figure(figsize=(7, 4))
sns.scatterplot(x="sepal_length", y="petal_length", hue="species", data=iris, palette="deep")
plt.title("Sepal Length vs Petal Length (Iris)")
plt.show()
```




7. Pairplot **Purpose:** Plots every pair of numerical columns against each other in one grid, with histograms on the diagonal. **When to use:** A fast, complete first look at relationships between *all* numeric features

```
sns.pairplot(iris, hue="species", palette="husl", diag_kind="kde")
plt.suptitle("Pairplot of Iris Dataset", y=1.02)
plt.show()
```

Pairplot of Iris Dataset



```
plt.figure(figsize=(6, 5))
sns.heatmap(iris.drop(columns="species").corr(), annot=True, cmap="YlGnBu")
plt.title("Correlation Heatmap — Iris Dataset")
plt.show()
```




Interpretation: petal_length and petal_width are almost perfectly correlated (0.96) — in a real project, you might drop one to reduce redundancy.

9. Line Plot

Purpose: Shows a trend of a numeric variable, typically over an ordered axis (time, sequence). **When to use:** Time-series data, or any variable naturally ordered along the x-axis.


```
bill_by_size = tips.groupby("size")["total_bill"].mean().reset_index()

plt.figure(figsize=(7, 4))
sns.lineplot(x="size", y="total_bill", data=bill_by_size, marker="o", color="darkgreen")
plt.title("Average Total Bill by Party Size (Tips Dataset)")
plt.xlabel("Party Size")
plt.ylabel("Average Total Bill ($)")
plt.show()
```




Interpretation: As party size increases, the average total bill rises steadily — larger groups naturally spend more, though the increase flattens for very large parties (fewer data points there too).

10. Pie Chart

Purpose: Shows the *proportion* each category contributes to the whole. **When to use:** A small number of categories (2–6) where showing "share of total" matters.

```
plt.figure(figsize=(5, 5))
titanic["embarked"].value_counts().plot(
    kind="pie", autopct="%1.1f%%", colors=sns.color_palette("Set3"), startangle=90)
plt.title("Passengers by Embarkation Port")
plt.ylabel("")
plt.show()
```




Interpretation: Over 72% of passengers embarked from Southampton ('S'), a much smaller share from Cherbourg ('C') and Queenstown ('Q').

Common mistake: Using pie charts for *many* categories (8+) — they become unreadable. A barplot is almost always a better choice beyond a handful of categories.

11. KDE Plot (Kernel Density Estimate)

Purpose: A smoothed, continuous version of a histogram — shows the estimated probability density of a variable. **When to use:** Comparing the *shape* of distributions across groups smoothly.


```
plt.figure(figsize=(7, 4))
sns.kdeplot(data=titanic, x="fare", hue="survived", fill=True, common_norm=False, palette="Set1")
plt.title("Fare Distribution by Survival (KDE Plot)")
plt.xlim(0, 300)
plt.show()
```



12. Strip Plot

Purpose: Plots every individual data point along a categorical axis (with slight jitter). **When to use:** Small-to-medium datasets, when you want to see *every* raw point, not just a summary.


```
plt.figure(figsize=(7, 4))
sns.stripplot(x="day", y="total_bill", data=tips, hue="day", jitter=True, palette="cool", legend=False)
plt.title("Total Bill by Day (Strip Plot)")
plt.show()
```



13. Swarm Plot

Purpose: Like a strip plot, but points are arranged to *never overlap* — shows density more clearly. **When to use:** Small datasets where you want an accurate, non-overlapping view of every point.

```
plt.figure(figsize=(7, 4))
sns.swarmplot(x="day", y="total_bill", data=tips, hue="day", palette="autumn", size=4, legend=False)
plt.title("Total Bill by Day (Swarm Plot)")
plt.show()
```




Interpretation: The 'swarm' shape reveals where bill amounts cluster for each day — e.g. most Thursday bills bunch around \$10–20, with few high outliers.

14. Jointplot

Purpose: Combines a scatter plot with histograms of both variables on the margins. **When to use:** Studying the relationship between two numeric variables *and* their individual distributions together.

```
sns.jointplot(x="total_bill", y="tip", data=tips, kind="reg", color="purple", height=6)
plt.suptitle("Total Bill vs Tip (Jointplot with Regression Line)", y=1.02)
plt.show()
```




Interpretation: There's a clear positive linear relationship — bigger bills lead to bigger tips. The side histograms show total_bill is right-skewed while tip is roughly normal.

15. Rug Plot

Purpose: Draws a small tick for every single data point along an axis — often layered under another plot. **When to use:** Adding raw data density detail underneath a histogram/KDE without extra clutter.

```
plt.figure(figsize=(7, 3))
sns.kdeplot(tips["tip"], color="teal")
sns.rugplot(tips["tip"], color="black")
```

```
plt.title("Tip Amount — KDE with Rug Plot")
plt.show()
```




Interpretation: The tick marks at the bottom show exactly where each real tip value falls — confirming most tips cluster between \$2–4, matching the KDE peak.

✓ 16. ECDF Plot (Empirical Cumulative Distribution Function)

Purpose: Shows what proportion of the data lies *at or below* each value. **When to use:** Comparing distributions precisely — answers "what % of values are under X?" directly.

```
plt.figure(figsize=(7, 4))
sns.ecdfplot(data=titanic, x="fare", hue="survived", palette="Dark2")
plt.title("ECDF of Fare by Survival")
plt.show()
```



Interpretation: At any fare value, you can read off the cumulative % of passengers at or below it. The non-survivors' curve rises faster at low fares — confirming they paid less on average than survivors.

Quick Check

- Which plot would you use to compare the *distribution shape* of `tip` between lunch and dinner?
- Which plot best shows the relationship between `total_bill` and `tip` together with their individual distributions?

Mini Assignment

Using the **Tips** dataset, create:

1. A boxplot of `total_bill` grouped by `time` (Lunch/Dinner).
2. A scatter plot of `total_bill` vs `tip`, colored by `smoker`.
3. Write 2 sentences interpreting each plot.

16. GroupBy & Aggregation

`groupby()` is one of the most powerful Pandas tools — it lets you **split** data into groups, **apply** a function to each group, and **combine** the results ("split-apply-combine").

```
# groupby() + agg() with multiple statistics at once
tips.groupby("day", observed=True)["total_bill"].agg(["mean", "median", "std", "count"])
```

	mean	median	std	count
day				
Thur	17.682742	16.20	7.886170	62
Fri	17.151579	15.38	8.302660	19
Sat	20.441379	18.24	9.480419	87
Sun	21.410000	19.63	8.832122	76

```
# pivot_table() – similar to groupby, but in a spreadsheet-like format
pd.pivot_table(tips, values="total_bill", index="day", columns="time",
               aggfunc="mean", observed=True)
```

time	Lunch	Dinner
day		
Thur	17.664754	18.780000
Fri	12.845714	19.663333
Sat	NaN	20.441379
Sun	NaN	21.410000

```
# crosstab() – frequency table between two categorical columns
pd.crosstab(titanic["pclass"], titanic["survived"], margins=True)
```

survived	0	1	All
pclass			
1	80	136	216
2	97	87	184
3	372	119	491
All	549	342	891

Interpretation: The `crosstab` above instantly shows survival counts per class — e.g. of 216 first-class passengers, 136 survived, confirming what we saw earlier in the barplot.

Tip: `groupby()` + `agg()` is your best friend for quickly answering business questions like *"What's the average bill per day?"* or *"Which class had the most deaths?"* without writing a single loop.

17. Full Data Cleaning Project — Titanic

Let's put **everything together** into one complete, real cleaning pipeline — exactly what you'd do before feeding this data into a Machine Learning model.

```
# Start fresh with a clean copy
clean_df = titanic.copy()
print("Starting shape:", clean_df.shape)
clean_df.head()
```

Starting shape: (891, 15)

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	who	adult_male	deck	embark_town	alive	alone
0	0	3	male	22.0	1	0	7.2500	S	Third	man	True	NaN	Southampton	no	False
1	1	1	female	38.0	1	0	71.2833	C	First	woman	False	C	Cherbourg	yes	False
2	1	3	female	26.0	0	0	7.9250	S	Third	woman	False	NaN	Southampton	yes	True
3	1	1	female	35.0	1	0	53.1000	S	First	woman	False	C	Southampton	yes	False
4	0	3	male	35.0	0	0	8.0500	S	Third	man	True	NaN	Southampton	no	True

Step 1 — Handle Missing Values

```
# 'age' -> fill with median (skewed, has outliers)
clean_df["age"] = clean_df["age"].fillna(clean_df["age"].median())

# 'embarked' -> fill with mode (categorical, very few missing)
clean_df["embarked"] = clean_df["embarked"].fillna(clean_df["embarked"].mode()[0])

# 'deck' -> too many missing (~77%) -> drop the column entirely
clean_df = clean_df.drop(columns=["deck"])

# drop rows still missing anything critical (safety net)
clean_df = clean_df.dropna(subset=["embark_town"])

print("Missing values remaining:\n", clean_df.isnull().sum()[clean_df.isnull().sum() > 0])
```

Missing values remaining:
Series([], dtype: int64)

Step 2 — Remove Duplicate Rows

```
before = clean_df.shape[0]
clean_df = clean_df.drop_duplicates()
after = clean_df.shape[0]
print(f"Removed {before - after} duplicate rows. New shape: {clean_df.shape}")
```

Removed 116 duplicate rows. New shape: (773, 14)

▼ Step 3 — Outlier Detection & Treatment (fare)

```
Q1, Q3 = clean_df["fare"].quantile([0.25, 0.75])
IQR = Q3 - Q1
lower, upper = Q1 - 1.5 * IQR, Q3 + 1.5 * IQR

n_outliers = clean_df[(clean_df["fare"] < lower) | (clean_df["fare"] > upper)].shape[0]
print(f"Fare outliers detected: {n_outliers}")

# Cap (winsorize) rather than delete, to avoid losing too many rows
clean_df["fare"] = clean_df["fare"].clip(lower=lower, upper=upper)
print("Fare range after capping:", clean_df["fare"].min(), "-", clean_df["fare"].max())
```

Fare outliers detected: 100
Fare range after capping: 0.0 - 71.675

▼ Step 4 — Encode Categorical Columns

```
# Drop redundant / high-cardinality / identifier-like columns first
clean_df = clean_df.drop(columns=["who", "adult_male", "alive", "class", "embark_town"])

# One-Hot Encode nominal columns
clean_df = pd.get_dummies(clean_df, columns=["sex", "embarked"], drop_first=True)

clean_df.head()
```

	survived	pclass	age	sibsp	parch	fare	alone	sex_male	embarked_Q	embarked_S
0	0	3	22.0	1	0	7.2500	False	True	False	True
1	1	1	38.0	1	0	71.2833	False	False	False	False
2	1	3	26.0	0	0	7.9250	True	False	False	True
3	1	1	35.0	1	0	53.1000	False	False	False	True
4	0	3	35.0	0	0	8.0500	True	True	False	True

Step 5 — Feature Scaling

```
num_cols = ["age", "fare"]
scaler = StandardScaler()
clean_df[num_cols] = scaler.fit_transform(clean_df[num_cols])

clean_df[num_cols].describe().round(2)
```

	age	fare
count	773.00	773.00
mean	0.00	-0.00
std	1.00	1.00
min	-2.12	-1.18
25%	-0.62	-0.82
50%	-0.11	-0.46
75%	0.47	0.33
max	3.68	2.04

Step 6 — Final Clean Dataset

```
print("Final cleaned shape:", clean_df.shape)
print("\nRemaining missing values:", clean_df.isnull().sum().sum())
clean_df.head()
```

Final cleaned shape: (773, 10)

Remaining missing values: 0

	survived	pclass	age	sibsp	parch	fare	alone	sex_male	embarked_Q	embarked_S
0	0	3	-0.548619	1	0	-0.851827	False	True	False	True
1	1	1	0.617360	1	0	2.019543	False	False	False	False
2	1	3	-0.257124	0	0	-0.821559	True	False	False	True
3	1	1	0.398739	1	0	1.204170	False	False	False	True
4	0	3	0.398739	0	0	-0.815954	True	True	False	True

```
# Save the final cleaned dataset to CSV
clean_df.to_csv("titanic_cleaned.csv", index=False)
print(" Saved as 'titanic_cleaned.csv'")
```

Saved as 'titanic_cleaned.csv'

Best practice recap: collect → inspect → handle missing values → remove duplicates → treat outliers → encode → scale → save. This exact order avoids most data leakage and preprocessing mistakes.

18. Mini Guided EDA Project — Tips Dataset

Now it's **your turn**! Follow the guided steps below, filling in the `# TODO` lines. A worked example is shown for the first step — try the rest yourself before checking Section 24's cheat sheet.

```
# STEP 1 (example, already solved): First look at the data
print(tips.shape)
tips.head()
```

```
(244, 7)
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

```
# STEP 2: Check for missing values
# TODO: use isnull().sum() on the tips dataset
tips.isnull().sum()
```

```

total_bill 0
tip         0
sex         0
smoker      0
day         0
time        0
size        0
dtype: int64
```


```
# STEP 3: Check for duplicate rows
# TODO: use duplicated().sum()
tips.duplicated().sum()
```

```
np.int64(1)
```

```
# STEP 4: Get summary statistics
# TODO: use describe()
tips.describe()
```

	total_bill	tip	size
count	244.000000	244.000000	244.000000
mean	19.785943	2.998279	2.569672
std	8.902412	1.383638	0.951100
min	3.070000	1.000000	1.000000
25%	13.347500	2.000000	2.000000
50%	17.795000	2.900000	2.000000
75%	24.127500	3.562500	3.000000
max	50.810000	10.000000	6.000000

```
# STEP 5: Visualize the relationship between total_bill and tip
# TODO: create a scatter plot, colored by 'time'
plt.figure(figsize=(7, 4))
sns.scatterplot(x="total_bill", y="tip", hue="time", data=tips, palette="Set1")
plt.title("Total Bill vs Tip, colored by Time of Day")
plt.show()
```



Now try on your own (no solution given!)

1. What is the average tip percentage ($\text{tip} / \text{total_bill} * 100$)? Add it as a new column.
2. Which day has the highest average tip percentage?
3. Do smokers tip more or less than non-smokers, on average?

19. Professional EDA Checklist

Use this checklist every single time you open a new dataset:

- ☐ **Load Data** — `pd.read_csv()` / `read_excel()` / `read_sql()`
- ☐ **Inspect** — `head()`, `tail()`, `sample()`, `shape`, `dtypes`
- ☐ **Summary** — `info()`, `describe()`, `describe(include='all')`
- ☐ **Missing Values** — `isnull().sum()`, heatmap, decide a fill/drop strategy
- ☐ **Duplicate Values** — `duplicated().sum()`, `drop_duplicates()`
- ☐ **Unique Values** — `unique()`, `nunique()`, `value_counts()`
- ☐ **Outliers** — boxplot, IQR rule, Z-score
- ☐ **Correlation** — correlation matrix + heatmap
- ☐ **Scaling** — `StandardScaler` / `MinMaxScaler` / `RobustScaler` (if needed)
- ☐ **Encoding** — Label / One-Hot / Ordinal (for categorical features)
- ☐ **Visualization** — histograms, boxplots, scatter plots, pairplots
- ☐ **Save Dataset** — export the cleaned data with `to_csv()`

Tip: Print this checklist out or pin it near your desk — following it in order saves hours of debugging later in the ML pipeline.

20. Practice Questions

Q1. Which function returns a dataset's dimensions (rows, columns)? A) `df.size` B) `df.shape` C) `df.ndim` D) `df.columns`

Q2. What is the key difference between `head()`, `tail()`, and `sample()`? A) They all return the same rows B) `head()` / `tail()` return fixed positions (start/end), `sample()` returns random rows C) `sample()` only works on numeric columns D) `tail()` shows the largest values

Q3. Why do we scale numerical data before training some ML models? A) To make the data prettier B) To ensure all features contribute fairly to distance/gradient-based algorithms C) Scaling is required for every ML algorithm without exception D) To remove missing values

Q4. Which pandas method shows the count of non-null values per column? A) `describe()` B) `info()` C) `shape` D) `columns`

Q5. Which value is most robust (least affected) by outliers? A) Mean B) Median C) Standard Deviation D) Range

- Q6. `nunique()` returns...** A) The unique values themselves B) The number of missing values C) The count of distinct values D) The most frequent value
- Q7. Which plot is best to detect outliers in a single numeric column?** A) Pie chart B) Boxplot C) Countplot D) Line plot
- Q8. What does IQR stand for and how is it calculated?** A) Inter Quartile Range = $Q3 - Q1$ B) Integer Quartile Ratio C) $Q1 + Q3$ D) Max - Min
- Q9. One-Hot Encoding is best suited for...** A) Ordinal categorical data with a natural rank B) Nominal categorical data with no natural order C) Purely numerical data D) Text paragraphs
- Q10. Which correlation value indicates the strongest negative relationship?** A) 0.95 B) 0.10 C) -0.85 D) 0.00
- Q11. `df.duplicated().sum()` tells you...** A) Total number of rows B) Number of duplicate rows C) Number of missing values D) Number of unique rows
- Q12. Which scaler is most robust to outliers?** A) MinMaxScaler B) StandardScaler C) RobustScaler D) Normalizer
- Q13. `groupby()` followed by `agg(['mean', 'count'])` will...** A) Return one aggregated statistic only B) Return multiple statistics per group in one call C) Remove all groups D) Only work on categorical columns
- Q14. A KDE plot is best described as...** A) A bar chart of frequencies B) A smoothed estimate of a variable's probability distribution C) A table of summary statistics D) A scatter plot with two variables
- Q15. Forward Fill (`ffill`) fills a missing value using...** A) The column mean B) A random value C) The previous valid value D) Zero
- Q16–Q25:** (For self-practice — write your own answer with one line of justification) 16. What's the difference between `.loc[]` and `.iloc[]`? 17. Why should you rarely use a pie chart for more than 6 categories? 18. What is the difference between Label Encoding and One-Hot Encoding? 19. Why is it risky to fill missing values using statistics from the full dataset before a train/test split? 20. What does a correlation of exactly 0 mean (and NOT mean)? 21. When would you prefer a median over a mean while filling missing values? 22. What is the difference between `value_counts()` and `unique()`? 23. Give one real-world example each of structured and unstructured data. 24. Why do tree-based models not require feature scaling? 25. What's the difference between variance and standard deviation?

► Click to reveal answers for Q1–Q15

Double-click (or enter) to edit

21. Programming Exercises

Complete each `# TODO` below yourself. Solutions are commented out — try first!

```
# Exercise 1: Load Iris and print the first five rows
# TODO: complete this
iris_ex = sns.load_dataset("iris")
print(iris_ex.head())
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
# Exercise 2: Find missing values in the Titanic dataset
# TODO: complete this
print(titanic.isnull().sum())
```


survived	0
pclass	0
sex	0
age	177
sibsp	0
parch	0
fare	0
embarked	2
class	0
who	0
adult_male	0
deck	688
embark_town	2
alive	0

```
alone          0  
dtype: int64
```


```
# Exercise 3: Remove duplicate rows from the Tips dataset (if any)  
# TODO: complete this  
tips_no_dup = tips.drop_duplicates()  
print("Before:", tips.shape, "After:", tips_no_dup.shape)
```

```
Before: (244, 7) After: (243, 7)
```

```
# Exercise 4: Plot a histogram of the 'total_bill' column from Tips  
# TODO: complete this  
plt.figure(figsize=(6, 4))  
sns.histplot(tips["total_bill"], bins=20, kde=True, color="coral")  
plt.title("Histogram of Total Bill")  
plt.show()
```



```
# Exercise 5: Create a correlation heatmap for the Penguins dataset  
# TODO: complete this  
plt.figure(figsize=(6, 5))  
sns.heatmap(penguins.select_dtypes(include=np.number).corr(), annot=True, cmap="coolwarm")  
plt.title("Penguins Correlation Heatmap")  
plt.show()
```



22. Mini Assignment — Full EDA on the Penguins Dataset

Perform a complete EDA on the `penguins` dataset (already loaded) and answer the following in new cells below. This is meant to be done **independently** — use everything you've learned in this notebook!

1. Number of rows and columns
2. Number of missing values (per column)
3. Number of duplicate rows
4. A correlation matrix / heatmap of the numeric columns
5. A boxplot showing `bill_length_mm` grouped by `species`
6. A pairplot colored by `species`
7. A fully cleaned version of the dataset (missing values handled, duplicates removed)

Use the empty code cell below to get started:


```
# Your Mini Assignment code goes here
penguins.head()
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex
0	Adelie	Torgersen	39.1	18.7	181.0	3750.0	Male
1	Adelie	Torgersen	39.5	17.4	186.0	3800.0	Female
2	Adelie	Torgersen	40.3	18.0	195.0	3250.0	Female
3	Adelie	Torgersen	NaN	NaN	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0	3450.0	Female

24. Final Summary — One-Page EDA Cheat Sheet

Load Dataset

`pd.read_csv()`, `sns.load_dataset()`





UNDERSTANDING ANY DATASET LIKE A DATA SCIENTIST

A Complete Guide to Exploratory Data Analysis (EDA) using Python Libraries

★ EXPLORE • UNDERSTAND • CLEAN • ANALYZE • VISUALIZE • INSIGHT • MODEL ★



WHY EDA?
..... EDA ROADMAP – STEP BY STEP
EDA CHECKLIST